


CS109 – Data Science

SVM, Performance evaluation

Joe Blitzstein, Hanspeter Pfister, Verena Kaynig-Fittkau




Announcements

- HW1 grades went out yesterday
- They are looking really good, well done everyone!
- HW2 is due this Thursday!
- You should submit an executed notebook
- But please without pages of test output


Recap K-NN

- Keeps all training data
- Training is fast
- Prediction is slow



Separating Hyperplane


- x : data point
- y : label $\in \{-1, +1\}$
- w : weight vector



$$w^T x = 0$$


Separating Hyperplane

- x : data point
- y : label $\in \{-1, +1\}$
- w : weight vector



Separating Hyperplane


- x : data point
- y : label $\in \{-1, +1\}$
- w : weight vector
- b : bias




$$w^T x + b = 0$$

Separating Hyperplane

- x : data point
- y : label $\in \{-1, +1\}$
- w : weight vector
- b : bias




Perceptron



$$w^T x + b = 0$$

Perceptron



Perceptron History

- invented 1957
- by Frank Rosenblatt
- the embryo of an electronic computer that [the Navy] expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence. (NYT 1958)



(<http://en.wikipedia.org/wiki/Perceptron>)



Perceptron.mp4

https://www.youtube.com/watch?v=cNxadbrN_al&list=PLdVOMWcqwwllaygvb9ZteZ1r4Br6kRuBO

Side Note: Step vs Sigmoid Activation




$$s(x) = \frac{1}{1 + e^{-cx}}$$


The Critics

- 1969: Minsky and Papert publish their book “Perceptrons”
- Very controversial book, some blame the book for causing the whole research area to stagnate.

The XOR Problem




The XOR Problem



Support Vector Machine

- Widely used for all sorts of classification problems
- Some people say it is the best of the shelf classifier out there

Maximum Margin Classification




Solution depends only on the support vectors!

Maximum Margin Classification

margin:

$$x_{\perp}^{(i)} = x^{(i)} - \gamma^{(i)} \cdot \frac{w}{\|w\|}$$

$$w^T x_{\perp}^{(i)} + b = 0$$

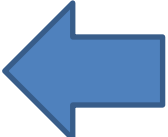


➔ $\gamma^{(i)} = \frac{w^T x^{(i)} + b}{\|w\|}$


Maximum Margin Classification

$$\gamma^{(i)} = y^{(i)}(w^T x + b)$$

$$\begin{array}{ll} \max_{\gamma, w, b} & \gamma \\ \text{s.t.} & y^{(i)}(w^T x^{(i)} + b) \geq \gamma, \quad i = 1, \dots, m \\ & ||w|| = 1. \end{array}$$

 non-convex



This Is Kind of Odd



- Which data points do we care the most about?
- What would those samples look like?

Á@ • ^ Á@ Á@ Á@ Á@ Á [• Á
 |ã ^ Á@ Á@ ! Á@ •

Two Very Similar Problems



What about outliers?

ξ_i : slack variables


$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2$$

$$\text{subject to:}$$



subject to:

$$y^{(i)}(w^T x^{(i)} + b) \geq 1 - \xi_i$$


$$(i = 1, \dots, n)$$




Two Very Similar Problems




Hard Margin ($C = \text{Infinity}$)



Soft Margin ($C = 10$)




XOR problem revised



Did we add information to make the problem separable?

Non-Linear Decision Boundary



SVM with a polynomial Kernel visualization

Created by:
Udi Aharoni

Quadratic Kernel

$$x = (x_1, x_2)$$

$$\Phi(x) = (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, x_2^2, \sqrt{2}x_1x_2)$$

$$\begin{aligned}\Phi(x) \cdot \Phi(z) &= 1 + 2 \sum_{i=1}^d x_i z_i \\ &\quad + \sum_{i=1}^d x_i^2 z_i^2 + 2 \sum_{i=1}^d \sum_{j=i+1}^d x_i x_j z_i z_j\end{aligned}$$

$$= (1 + x \cdot z)^2$$

Kernel Functions

$$K(x, z) = \Phi(x) \cdot \Phi(z)$$

- Polynomial:

$$K(x, z) = (1 + x \cdot z)^s$$

- Radial basis function (RBF):

$$K(x, z) = \exp(-\gamma(x - z)^2)$$


So what is the excitement?

$$\max_{\alpha} \sum$$

$$\text{s.t. } \alpha_i$$

$$\sum$$

$$(i)^T x(j)$$



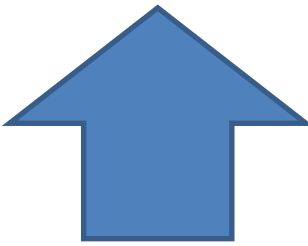
$$\arg \min$$

$$\text{s.t. } y$$


"I THINK YOU SHOULD BE MORE EXPLICIT
HERE IN STEP TWO."

So what is the excitement?

$$\begin{aligned}
 & \max_{\alpha} \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \boxed{x^{(i)T} x^{(j)}} \\
 & \text{s.t. } \alpha_i \geq 0, \quad i = 1, \dots, m \\
 & \quad \sum_{i=1}^m \alpha_i y^{(i)} = 0
 \end{aligned}$$



$$\begin{aligned}
 & \arg \min_{w,b} \frac{1}{2} ||w||^2 \\
 & \text{s.t. } y^{(i)} (w^T x^{(i)} + b) \geq 1
 \end{aligned}$$



 $\boxed{K(x^{(i)}, x^{(j)})}$

Prediction

$$w^T x + b = \sum_{i=1}^m \alpha_i y^{(i)} \langle x^{(i)}, x \rangle + b.$$

- Again we can use the kernel trick!
- Prediction speed depends on number of support vectors

The Miracle Explained

- Andrew Ng does this really well
- <http://cs229.stanford.edu/notes/cs229-notes3.pdf>
- Course is also on Youtube, ItunesU, etc.

Kernel Trick for SVMs

- Arbitrary many dimensions
- Little computational cost
- Maximal margin helps with curse of dimensionality

Face Recognition

pred: Colin_Powell
true: Colin_Powell



pred: George_W_Bush
true: George_W_Bush



pred: Colin_Powell
true: Colin_Powell



pred: Tony_Blair
true: Tony_Blair



pred: George_W_Bush
true: George_W_Bush



pred: Colin_Powell
true: Colin_Powell



pred: George_W_Bush
true: George_W_Bush



pred: George_W_Bush
true: George_W_Bush



pred: Tony_Blair
true: Tony_Blair



pred: Colin_Powell
true: Colin_Powell



pred: George_W_Bush
true: George_W_Bush



pred: Donald_Rumsfeld
true: Donald_Rumsfeld



Face Recognition

- Load image data
 - Put your test data aside
 - Extract Eigenfaces
 - Train SVM
 - Evaluate performance
-
- Red are cross validation steps

http://scikit-learn.org/stable/auto_examples/applications/face_recognition.html#example-applications-face-recognition-py






SVM_sign_language.mp4




[Jhon Gonzalez](#)

https://www.youtube.com/watch?v=cxHMgl2_5zg




$\gamma=10^{-1}, C=10^{-2}$ $\gamma=10^0, C=10^{-2}$ $\gamma=10^1, C=10^{-2}$



$\gamma=10^{-1}, C=10^0$ $\gamma=10^0, C=10^0$ $\gamma=10^1, C=10^0$



$\gamma=10^{-1}, C=10^2$ $\gamma=10^0, C=10^2$ $\gamma=10^1, C=10^2$



Tips and Tricks

- SVMs are not scale invariant
- Check if your library normalizes by default
- Normalize your data
 - mean: 0 , std: 1
 - map to $[0,1]$ or $[-1,1]$
- Normalize test set in same way!

Tips and Tricks

- RBF kernel is a good default
- For parameters try exponential sequences
- Read:


Chih-Wei Hsu et al., “**A Practical Guide to Support Vector Classification**”,
Bioinformatics (2010)

SVM vs KNN

- What are the main key differences?


Parameter Tuning

- Given a classification task
- Which kernel ?
- Which kernel parameter values?
- Which value for C?




Try different combinations
and take the **best**.

Train vs. Test Error



Where is KNN on this graph for $K=1$, or for $K=\text{Inf}$?

Grid Search



Zang et al., "Identification of heparin samples that contain impurities or contaminants by chemometric pattern recognition analysis of proton NMR spectral data", Anal Bioanal Chem (2011)

Error Measures

- True positive (tp)
- True negative (tn)
- False positive (fp)
- False negative (fn)

		predicted	
		1	-1
true	1	tp	fn
	-1	fp	tn

TPR and FPR

- True Positive Rate:

$$\frac{tp}{tp + fn}$$

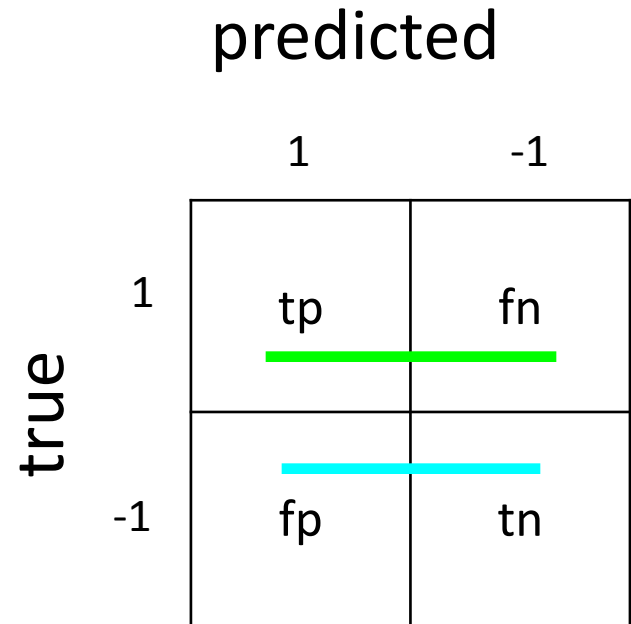
- False Positive Rate:

$$\frac{fp}{fp + tn}$$


predicted

	1	-1
1	tp	fn
-1	fp	tn


true

A confusion matrix diagram. The vertical axis is labeled 'true' with values 1 and -1. The horizontal axis is labeled 'predicted' with values 1 and -1. The matrix cells are labeled: top-left is 'tp', top-right is 'fn', bottom-left is 'fp', and bottom-right is 'tn'. A thick green horizontal bar is drawn across the 'tp' and 'fn' cells. A thick cyan horizontal bar is drawn across the 'fp' and 'tn' cells.

Receiver Operating Characteristic



ROC Example



Precision Recall

- Recall: $\frac{tp}{tp + fn}$

- Precision: $\frac{tp}{tp + fp}$

predicted


	1	-1
1	tp	fn
-1	fp	tn

true

Precision Recall



- **Recall:** If I pick a random positive example, what is the probability of making the right prediction?
- **Precision:** If I take a positive prediction example, what is the probability that it is indeed a positive example?

Precision Recall Curve



Comparison

Algorithm 1: $\frac{1}{2} \left(\frac{1}{2} + \frac{1}{2} \right)$
Algorithm 2: $\frac{1}{2} \left(\frac{1}{2} + \frac{1}{2} \right)$



J. Davis & M. Goadrich,
“The Relationship Between Precision-Recall and ROC Curves.”,
ICML (2006)

F-measure


- Weighted average of precision and recall

$$F_{\beta} = \frac{(\beta^2 + 1) \cdot P \cdot R}{\beta^2 \cdot P + R}$$

- Usual case: $\beta = 1$
- Increasing β allocates weight to recall

Multi Class

À^|||, À•Àaà*![") à&|æ•QMA^àEÀ|^E
À^à&•Àaà*![") à&|æ•QMA^|||, ÀEÀ|^D




One vs. All

One vs All

- Train n classifier for n classes
- Take classification with greatest margin
- Slow training

Multi Class


$\vec{A} \in \mathbb{R}^n$, $\vec{B} \in \mathbb{R}^m$, $\vec{C} \in \mathbb{R}^k$, $\vec{D} \in \mathbb{R}^l$



One vs One

- Train $n(n-1)/2$ classifiers
- Take majority vote
- Fast training

Confusion Matrix



Recap

- Perceptrons are great
- But really just a separating hyperplane
- So is SVM
- Kernels are neat
- Evaluation metrics are important